Cristian Hohbein

Brea, CA 92823 | chohbein2@gmail.com | +1 (714) 510-1314 | LinkedIn | chohbein.github.io

PROFESSIONAL SUMMARY

Engineer who builds and runs data pipelines end-to-end: cloud ingestion, scheduled transforms, and the production services on top of them. Sole engineer on a live data stack, which means also being the one who fixes it when a job breaks. Most recent work has been on LLM-powered applications and how to evaluate them. M.S. Data Science, Fall 2026.

EDUCATION

University of Colorado, Boulder M.S. Data Science | Fall 2026

Relevant Coursework: Big Data Analytics, Supervised/Unsupervised Algorithms in Machine Learning, Deep Learning, Modern AI Models for Vision and MultiModal Understanding

University of California, Santa Barbara B.A. Data Science | June 2024

Relevant Coursework: Advanced Statistical Models, Big Data Analytics, Statistical Machine Learning

EXPERIENCE

Co-Founder & Engineer (Data / ML) 2025 – Present

Cartez

- Built and operate the full data pipeline behind a live production app, as sole engineer: ingestion and transformation of 340M+ user interactions and 2.6M items into modeled, queryable outputs in PostgreSQL.
- Run scheduled daily batch refreshes through a Supabase edge function, using pool rotation and caching to keep the job reliable and hold down compute cost.
- Deployed and configured the supporting services behind the pipeline: a FastAPI service on Render backed by pgvector similarity search, tuned for low-latency reads.
- Own the whole stack end-to-end, from debugging data and job failures in production to designing the models and abstractions other parts of the product are built on.

PROJECTS

News Aggregator (Scheduled ETL on AWS) github.com/chohbein/newsapp

- A site that collects news from across the web and groups related coverage into topics. Built the scheduled ETL pipeline behind it on AWS (Lambda, S3, RDS), ingesting 100+ articles daily and transforming unstructured text into structured records in PostgreSQL.
- Added downstream enrichment with Sentence-BERT embeddings and DBSCAN clustering, plus keyword and summary extraction, turning raw ingested text into datasets the application could query directly.

Tools: AWS Lambda, S3, RDS, PostgreSQL, Python, Sentence-BERT, spaCy, DBSCAN

LLM-Powered Language Feedback API github.com/chohbein/language-feedback-api

- An API that returns grammar corrections and difficulty ratings for language learners. Built as a Dockerized FastAPI service returning schema-conformant structured output via Claude tool use, with Pydantic validation so downstream consumers can trust the shape of what comes back.
- Defined success metrics and built a pytest evaluation suite covering error detection, span precision, and rating accuracy, used to surface and measure LLM failure modes rather than assume correctness.
- Reduced latency and API cost with request-level caching keyed on sentence and language pair.

Tools: FastAPI, Pydantic, pytest, Docker, prompt engineering, tool use, Anthropic Claude API

Esports Match Prediction Pipeline

- Predicts winners of professional esports matches from team and player history. Built the whole pipeline from scratch: scraped and normalized historical match statistics, engineered features, and trained an XGBoost model with Optuna hyperparameter search.
- Traced a persistent discrepancy between validation and live results to its root cause rather than patching the symptom: the test set had been silently steering hyperparameter search and early stopping. Rebuilt the evaluation around a single-pass held-out test, and the corrected metrics matched the predicted size of the error.

Tools: Python, XGBoost, Optuna, Pandas, NumPy, web scraping

SKILLS

Data & Pipelines: SQL, PostgreSQL, Python, Pandas, NumPy, PySpark (Apache Spark), ETL design, scheduled batch jobs, data modeling, pgvector
Cloud & Tooling: AWS (Lambda, S3, RDS, EC2, SageMaker), Docker, Git, FastAPI, pytest, Supabase
AI & LLMs: Anthropic Claude API, prompt engineering, tool use / structured output, LLM evaluation suites, embeddings, Sentence-BERT, semantic search
Analytics & ML: scikit-learn, XGBoost, PyTorch, Power BI, Matplotlib, Seaborn, model evaluation & validation, clustering, recommender systems